

NAIRR Pilot

National Artificial Intelligence
Research Resource Pilot

Workflows and AI

Mats Rynge, USC Information Sciences Institute / NSF ACCESS

April 3 / Track 1

AI Workshop Denver, CO April 2-3, 2025



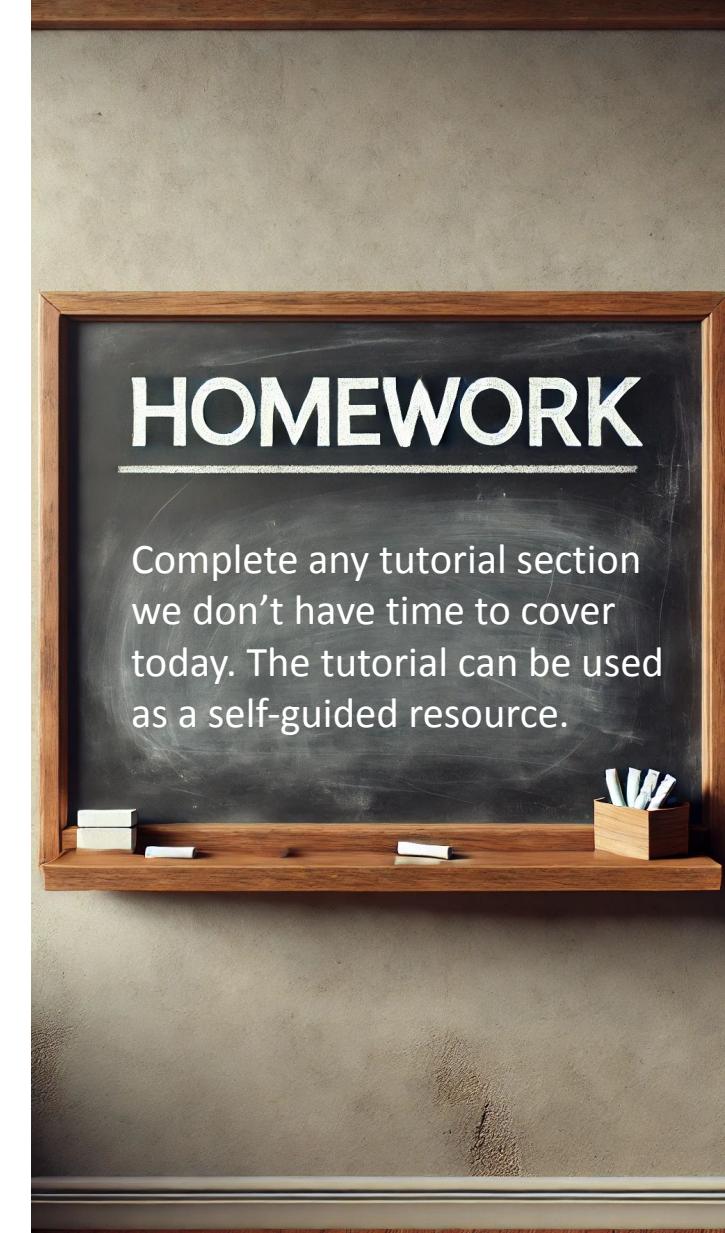
This Session

Main takeaways:

- 1. Use a workflow system for your AI workloads**
- 2. Scheduler overlay to simplify resource management**

Why ACCESS Pegasus?

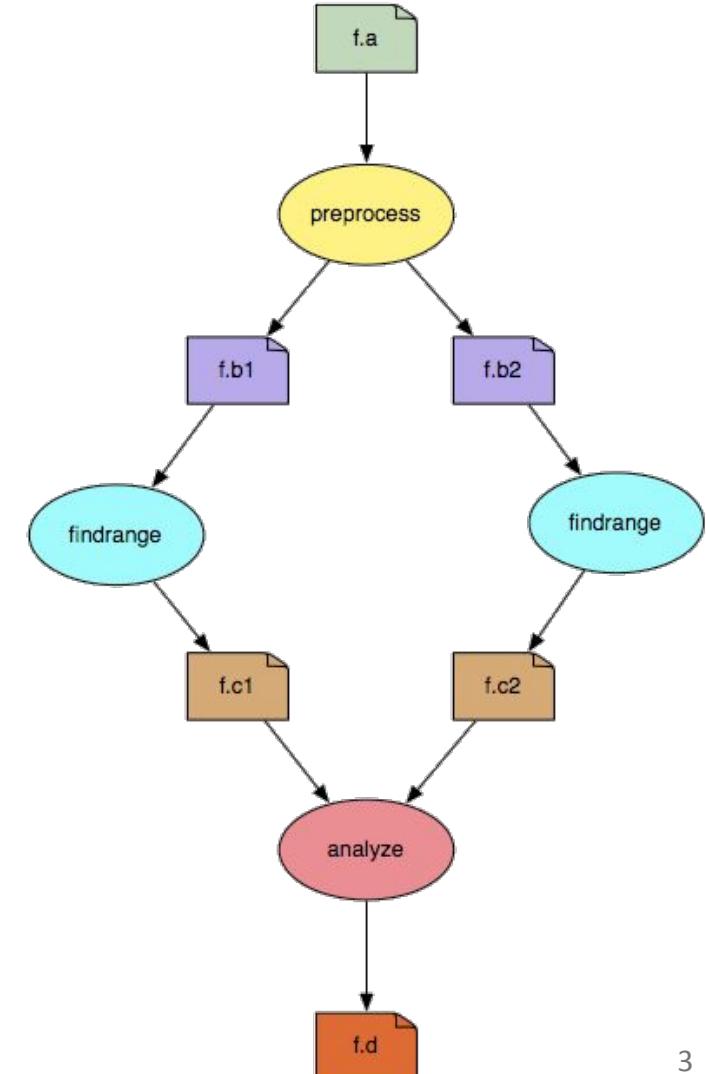
Pegasus can be deployed and used in a variety of different ways, such as cli/Jupyter, local install on your machine, a cluster or cloud. One option is a hosted system like the ACCESS Pegasus one we are using today.





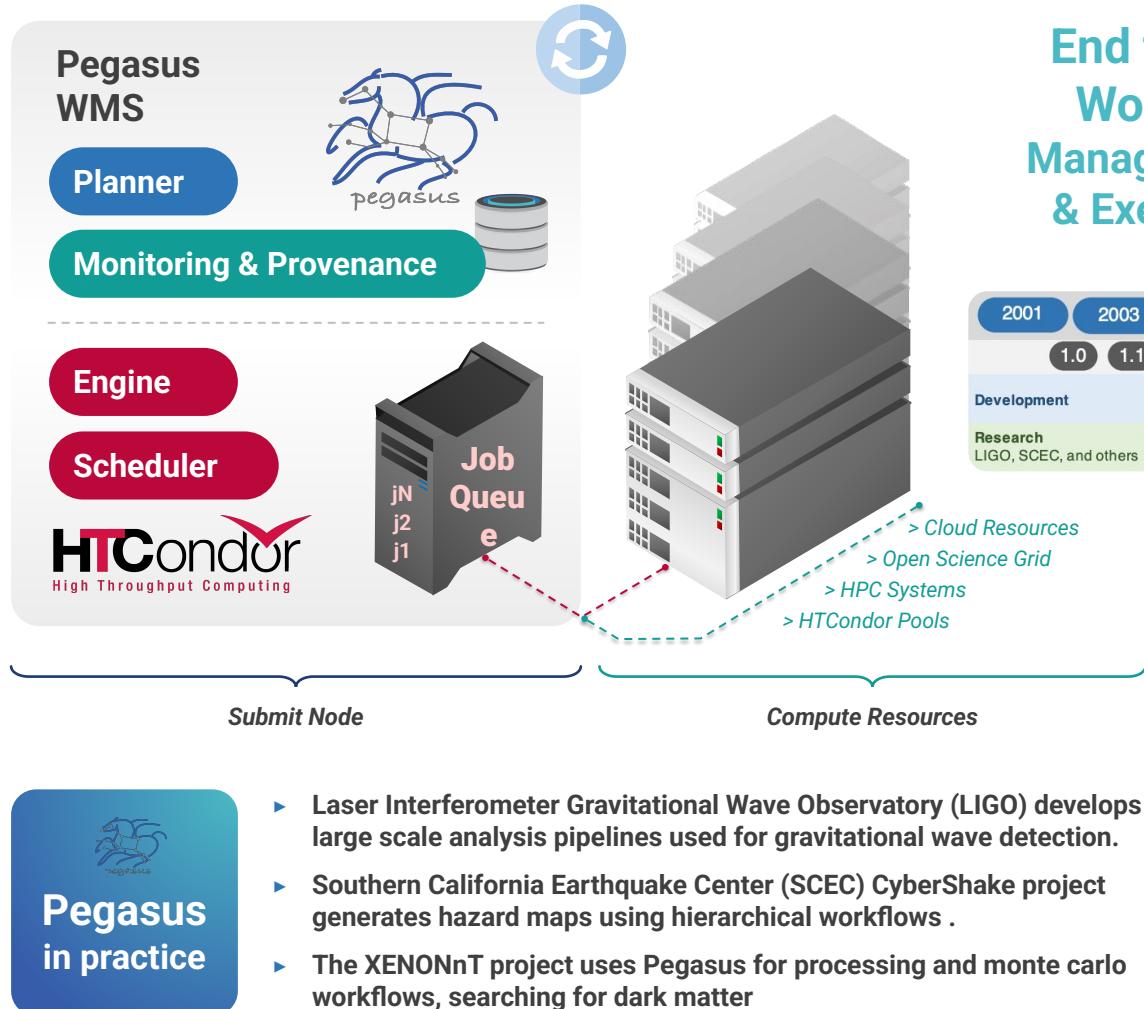
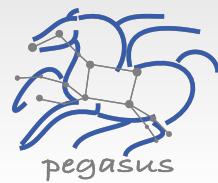
Scientific Workflows

- An abstraction to express ensemble of complex computational operations
 - *Eg: retrieving data from remote storage services, executing applications, and transferring data products to designated storage sites*
- A workflow is represented as a directed acyclic graph (DAG)
 - *Nodes: tasks or jobs to be executed*
 - *Edges: depend between the tasks*
- Have a monolithic application/experiment?
 - *Find the inherent DAG structure in your application to convert into a workflow*



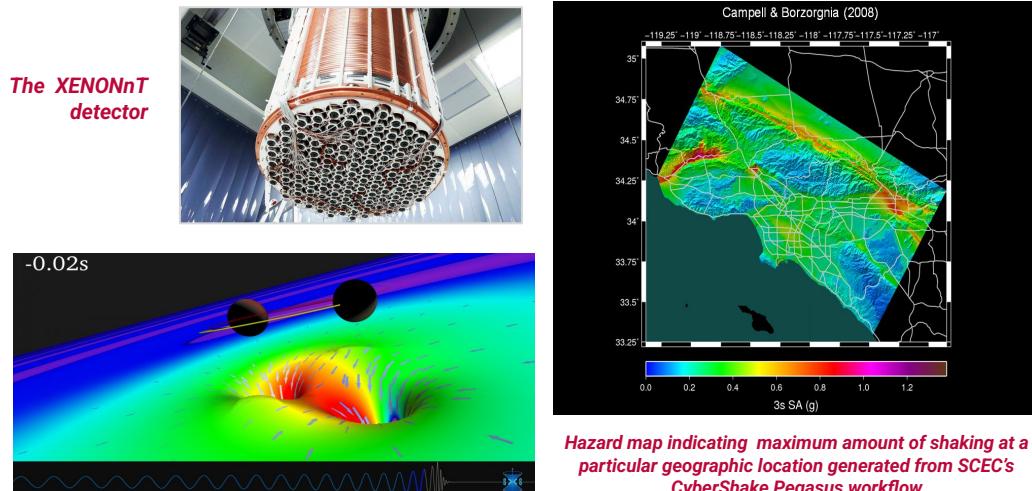


Pegasus Workflow Management System



End to End Workflow Management & Execution

- ▶ Develop portable scientific workflows in Python, Java, and R
- ▶ Compile workflows to be run on heterogeneous resources
- ▶ Monitor and debug workflow execution via CLI and web-based tools
- ▶ Recover from failures with built-in fault tolerance mechanisms
- ▶ Regular release schedule incorporating latest research and development



- ▶ Laser Interferometer Gravitational Wave Observatory (LIGO) develops large scale analysis pipelines used for gravitational wave detection.
- ▶ Southern California Earthquake Center (SCEC) CyberShake project generates hazard maps using hierarchical workflows .
- ▶ The XENONnT project uses Pegasus for processing and monte carlo workflows, searching for dark matter

Pegasus in practice

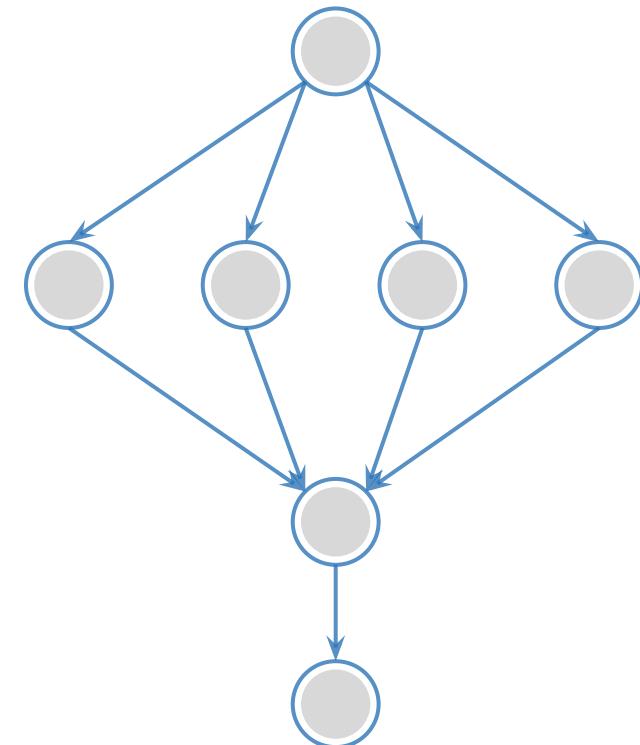


▲ **Pegasus WMS == Pegasus planner (mapper) + DAGMan workflow engine + HTCondor scheduler/broker**

- Pegasus maps workflows to infrastructure
- DAGMan manages dependencies and reliability
- HTCondor is used as a broker to interface with different schedulers

▲ **Workflows are DAGs**

- Nodes: jobs, edges: dependencies
- No while loops, no conditional branches
- Jobs are standalone executables



▲ **Planning occurs ahead of execution**

▲ **Planning converts an abstract workflow into a concrete, executable workflow**

- Planner is like a compiler



▲ **Pegasus WMS == Pegasus planner (mapper) + DAGMan workflow engine + HTCondor scheduler/broker**

- Pegasus maps workflows to infrastructure
- DAGMan manages dependencies and reliability
- HTCondor is used as a broker to interface with different schedulers

▲ **Workflows are DAGs**

- Nodes: jobs, edges: dependencies
- No while loops, no conditional branches
- Jobs are standalone executables

▲ **Planning occurs ahead of execution**

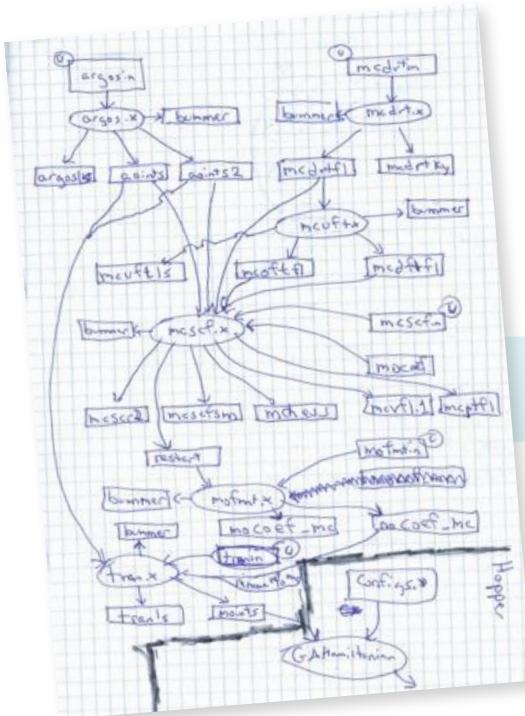
▲ **Planning converts an abstract workflow into a concrete, executable workflow**

- Planner is like a compiler

Pegasus reasons about data, in order to do automatic placement, movement, and integrity.



Pegasus provides APIs
to generate the Abstract Workflow



```
#!/usr/bin/env python3

import os
import logging
from pathlib import Path
from argparse import ArgumentParser

logging.basicConfig(level=logging.DEBUG)

# --- Import Pegasus API -----
from Pegasus.api import *

# --- Create Abstract Workflow -----
wf = Workflow("pipeline")

webpage = File("pegasus.html")

# --- Create Parent Job -----
curl_job = (
    Job("curl")
    .add_args("-o", webpage, "http://pegasus.isi.edu")
    .add_outputs(webpage, stage_out=False, register_replica=False)
)

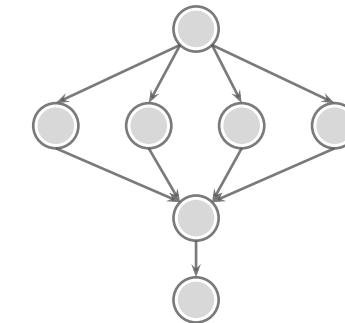
count = File("count.txt")

# --- Create Dependent Job -----
wc_job = (
    Job("wc")
    .add_args("-l", webpage)
    .add_inputs(webpage)
    .set_stdout(count, stage_out=True, register_replica=True)
)

# --- Add jobs to the Abstract Workflow -----
wf.add_jobs(curl_job, wc_job)

# --- Add control flow dependency -----
wf.add_dependency(wc_job, parents=[curl_job])

# --- Write out the Abstract Workflow -----
wf.write()
```



Abstract Workflow

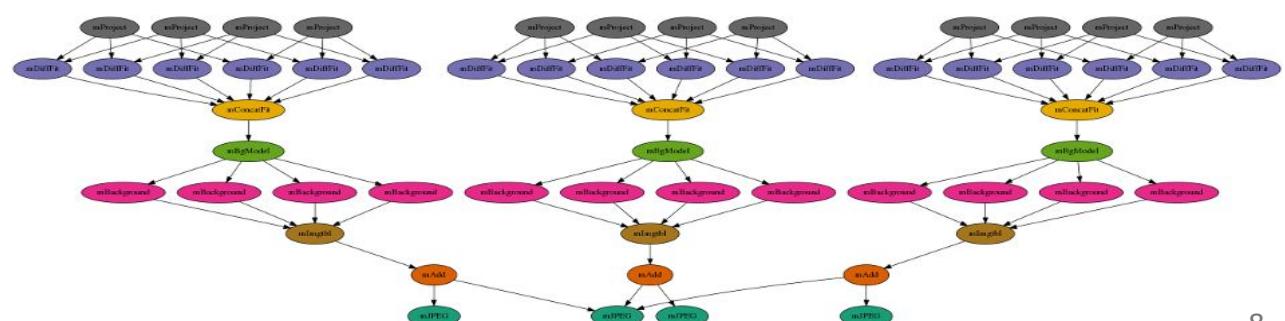
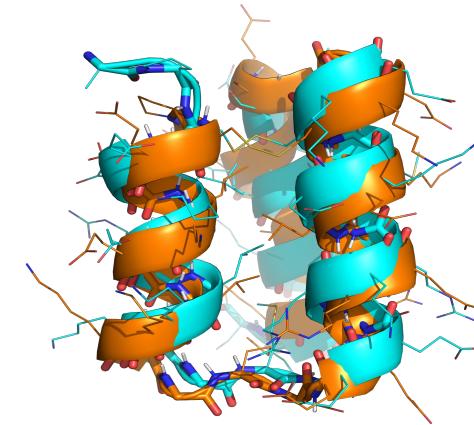
```
x-pegasus:  
  apilang: python  
  createdBy: vahi  
  createdOn: 11-19-20T14:57:58Z  
pegasus: '5.0'  
name: pipeline  
jobs:  
- type: job  
  name: curl  
  id: ID00000001  
  arguments:  
  - -o  
  - pegasus.html  
  - http://pegasus.isi.edu  
  uses:  
  - lfn: pegasus.html  
    type: output  
    stageOut: false  
    registerReplica: false  
- type: job  
  name: wc  
  id: ID00000002  
  stdout: count.txt  
  arguments:  
  - -l  
  - pegasus.html  
  uses:  
  - lfn: count.txt  
    type: output  
    stageOut: true  
    registerReplica: true  
  - lfn: pegasus.html  
    type: input  
jobDependencies:  
- id: ID00000001  
  children:  
  - ID00000002
```



Example Workflows

In addition to tutorial workflows, a set of example workflows are automatically installed into each user account - easy to explore, execute and modify!

- Artificial Intelligence
 - Lung Segmentation
 - Mask Detection
 - Orca Sound
 - LLM + RAG
- Astronomy
 - Montage
- Bioinformatics
 - AlphaFold
 - Rosetta
 - VariantCalling





LLM - RAG

LLM RAG (Large Language Model Retrieval-Augmented Generation) is a technique that enhances large language models by incorporating information retrieval mechanisms.

It involves retrieving relevant information from a database or document corpus and combining it with the original query to provide additional context. This augmented input is then processed by the large language model to generate more accurate and contextually relevant responses.

LLM RAG offers benefits such as improved accuracy, access to up-to-date information, and better contextual understanding, making it useful for applications like question answering, summarization, and conversational AI.

1. `gpu96`: This flavor provides 4 NVIDIA V100 GPUs (with 32GB of memory each) and 128GB of RAM on the host CPU node.
2. `gpu60`: This flavor offers 2 NVIDIA V100 GPUs (each with 32GB of memory) and 64GB of RAM on the host CPU node.

The Jetstream2 system offers several GPU instance flavors: g3.small, g3.medium, g3.large, and g3.xl. Here's a summary of their specifications:

- g3.small: 4 vCPUs, 15 GB RAM, 60 GB local storage, 20% GPU compute, 20 GB GPU RAM
- g3.medium: 8 vCPUs, 30 GB RAM, 60 GB local storage, 25% GPU compute, 10 GB GPU RAM
- g3.large: 16 vCPUs, 60 GB RAM, 60 GB local storage, 50% GPU compute, 20 GB GPU RAM
- g3.xl: 32 vCPUs, 125 GB RAM, 60 GB local storage, 100% GPU compute, 40 GB GPU RAM



Tutorial: LLM - RAG

Goal: Execute a set of LLMs, locally on NAIRR resources, and use RAG to augment the models and answer a set of predefined prompts.



Container includes full LLM and toolchain (Mistral, LangChain, Chroma)

Workflow Input: Public domain books (Project Gutenberg)

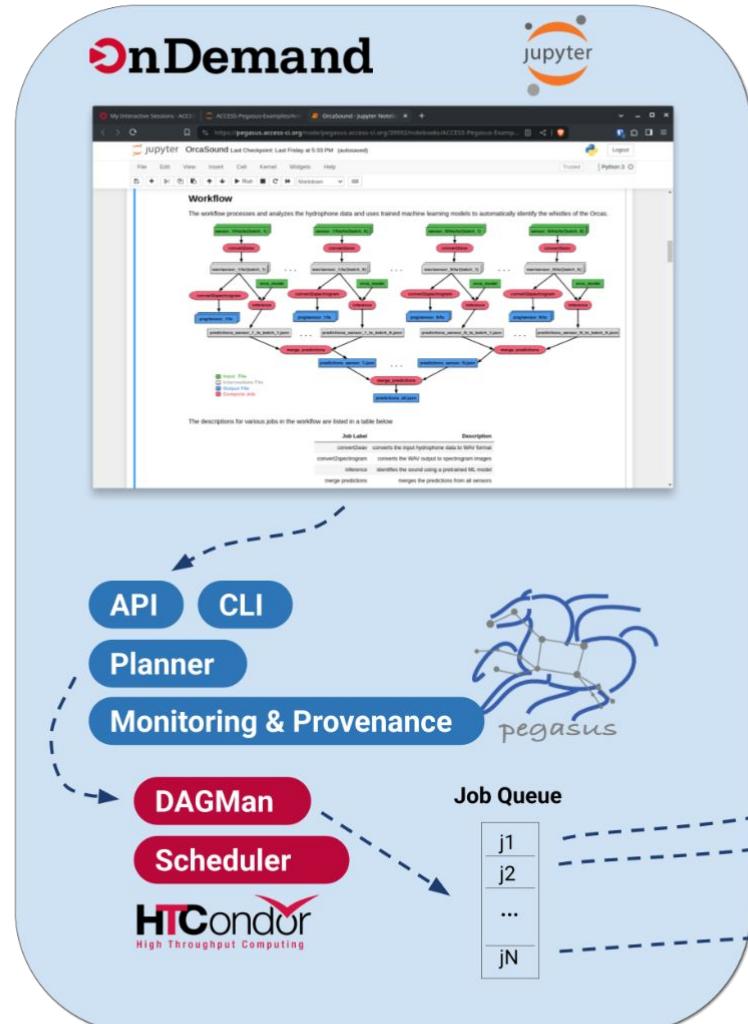
Prompts:

- Please provide a one paragraph summary of the book.
- Who is the protagonist in the book?
- Who is the antagonist in the book?
- What time period is the book set in?

Can execute on GPUS from any of the capacity providers: TestPool, Cloud, HTCondor Annex, OSPool



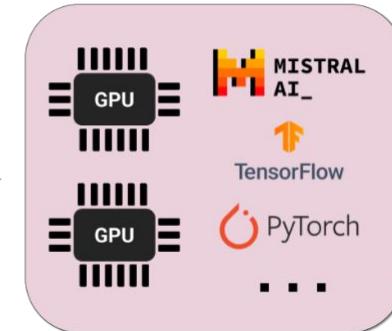
Chroma



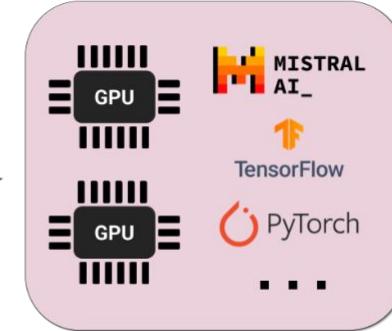
pegasus.access-ci.org



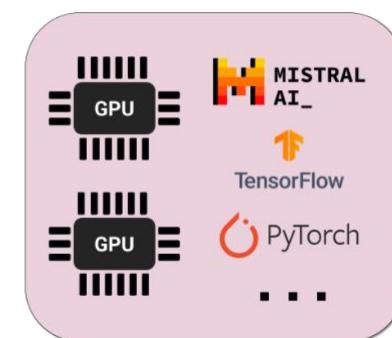
remote job submission



TestPool



Cloud



HPC

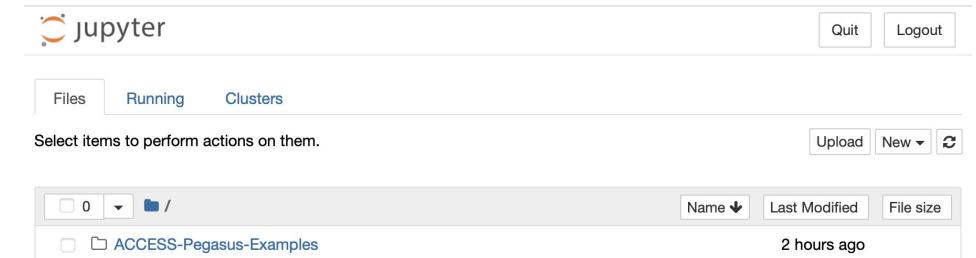


Tutorial Setup

Set of Jupyter Notebooks, contains a mix of overview, pointers, workflows. Step through them one cell at the time, starting from the top.

Log in to ACCESS Pegasus, and start a Jupyter Notebook.

Tutorials can be found in the ***ACCESS-Pegasus-Examples*** directory.



Using your ACCESS account, verify that you can use a web browser to log into:

<https://pegasus.access-ci.org>

NAIRR Pilot

National Artificial Intelligence
Research Resource Pilot

Hands on: Running our first workflow

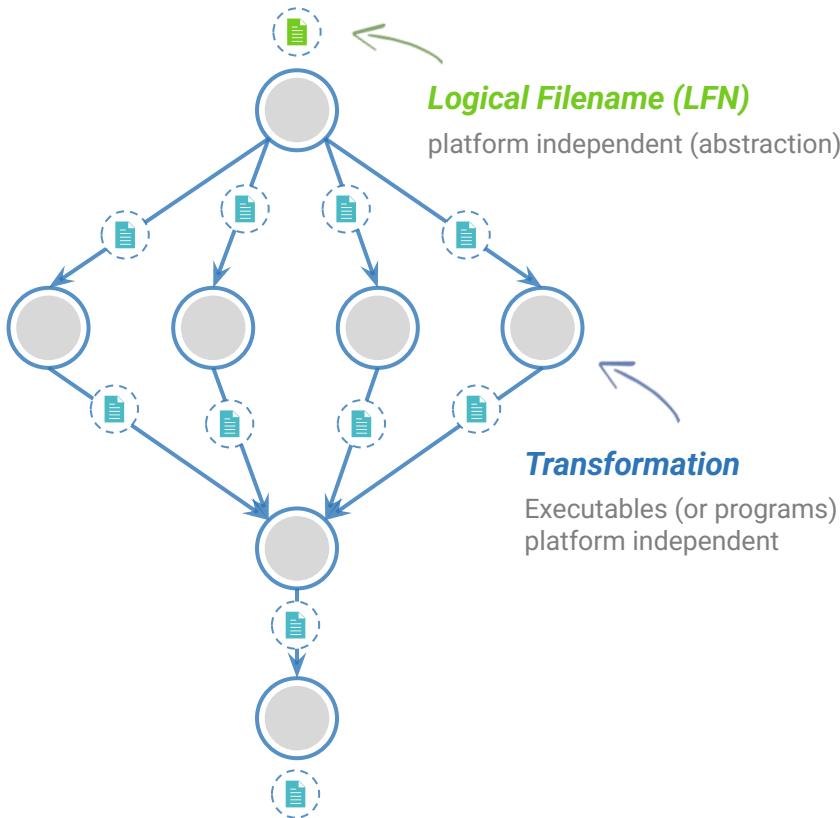
<https://pegasus.access-ci.org>

Input Workflow Specification YAML formatted

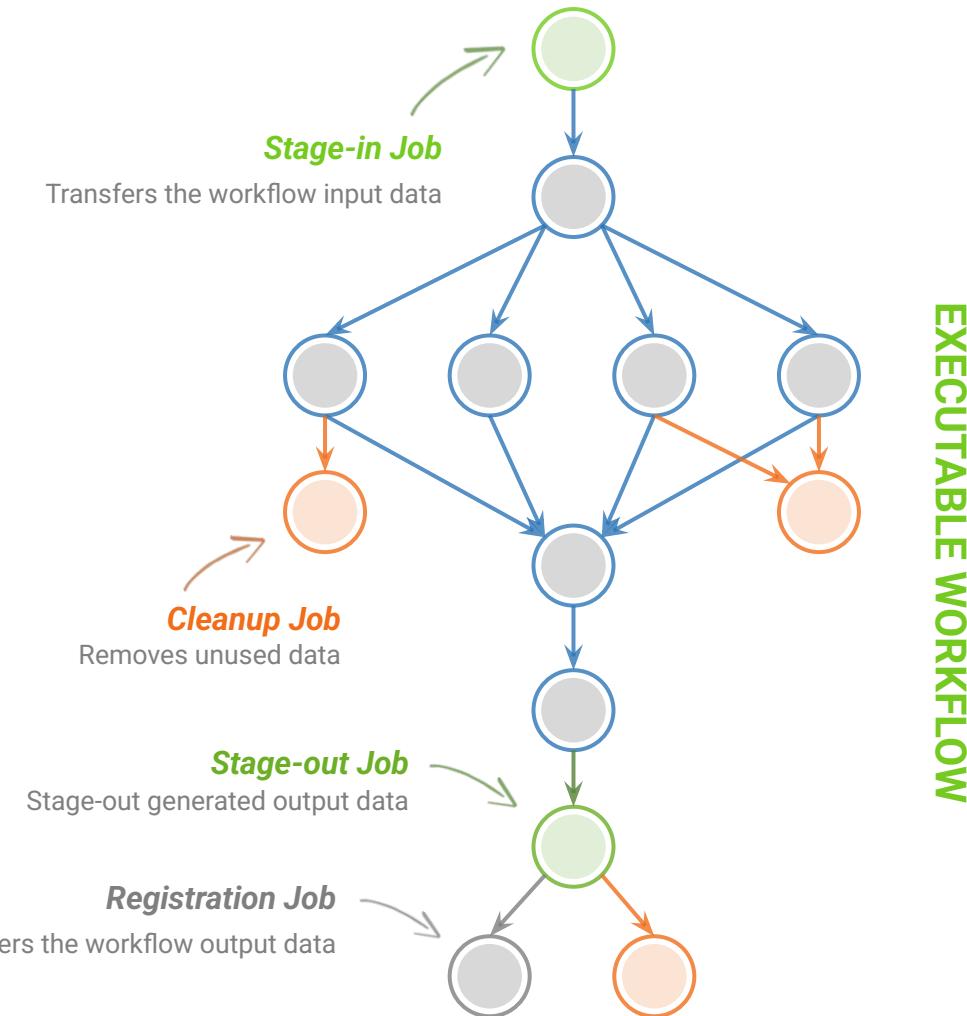
Portable Description

Users do not worry about low level execution details

ABSTRACT WORKFLOW



Output Workflow



EXECUTABLE WORKFLOW



So, what other information does Pegasus need?

Site Catalog

Describes the sites where
The workflow jobs are to be executed

Transformation Catalog

Describes all of the executables
(called “transformations”)
used by the workflow

Replica Catalog

Describes all of the input data stored
on external servers



NAIRR Pilot

National Artificial Intelligence
Research Resource Pilot

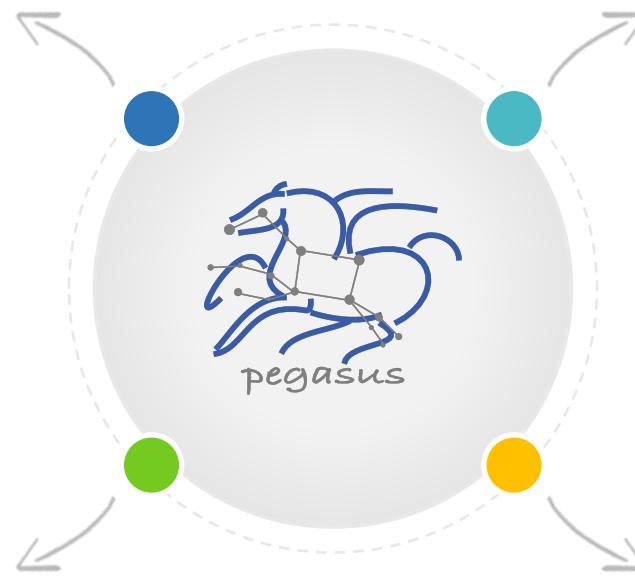
Hands on: Catalogs

And if a job fails?



Postscript

detects non-zero exit code output
parsing for success or failure
message exceeded timeout do not
produced expected output files



Job Retry



helps with transient failures
set number of retries per job
and run



Checkpoint Files

job generates checkpoint files
staging of checkpoint files is
automatic on restarts

Rescue DAGs



workflow can be restarted from
checkpoint file recover from
failures with minimal loss

command-line...

```
$ pegasus-status pegasus/examples/split/run0001
STAT IN_STATE JOB
Run 00:39 split-0 (/home/pegasus/examples/split/run0001)
Idle 00:03 └split_ID0000001
Summary: 2 Condor jobs total (I:1 R:1)

UNRDY READY PRE IN_Q POST DONE FAIL %DONE STATE DAGNAME
 14      0     0    1      0    2     0    11.8 Running *split-0.dag
```

```
$ pegasus-analyzer pegasus/examples/split/run0001
pegasus-analyzer: initializing...

*****Summary*****
Total jobs : 7 (100.00%)
# jobs succeeded : 7 (100.00%)
# jobs failed : 0 (0.00%)
# jobs unsubmitted : 0 (0.00%)
```

```
$ pegasus-statistics -s all pegasus/examples/split/run0001
-----
Type      Succeeded Failed Incomplete Total Retries Total+Retries
Tasks        5       0       0       5       0       5
Jobs         17      0       0      17      0      17
Sub-Workflows  0       0       0       0       0       0
-----
Workflow wall time : 2 mins, 6 secs
Workflow cumulative job wall time : 38 secs
Cumulative job wall time as seen from submit side : 42 secs
Workflow cumulative job badput wall time :
Cumulative job badput wall time as seen from submit side :
```

**Provenance Data
can be Summarized
*pegasus-statistics***
or
**Used for Debugging
*pegasus-analyzer***

NAIRR Pilot

National Artificial Intelligence
Research Resource Pilot

Hands on: Statistics

NAIRR Pilot

National Artificial Intelligence
Research Resource Pilot

Hands on: Statistics and Debugging

NAIRR Pilot

National Artificial Intelligence
Research Resource Pilot

Resources Overlay



Advantages of Glidein (Pilot) Jobs in HTCondor



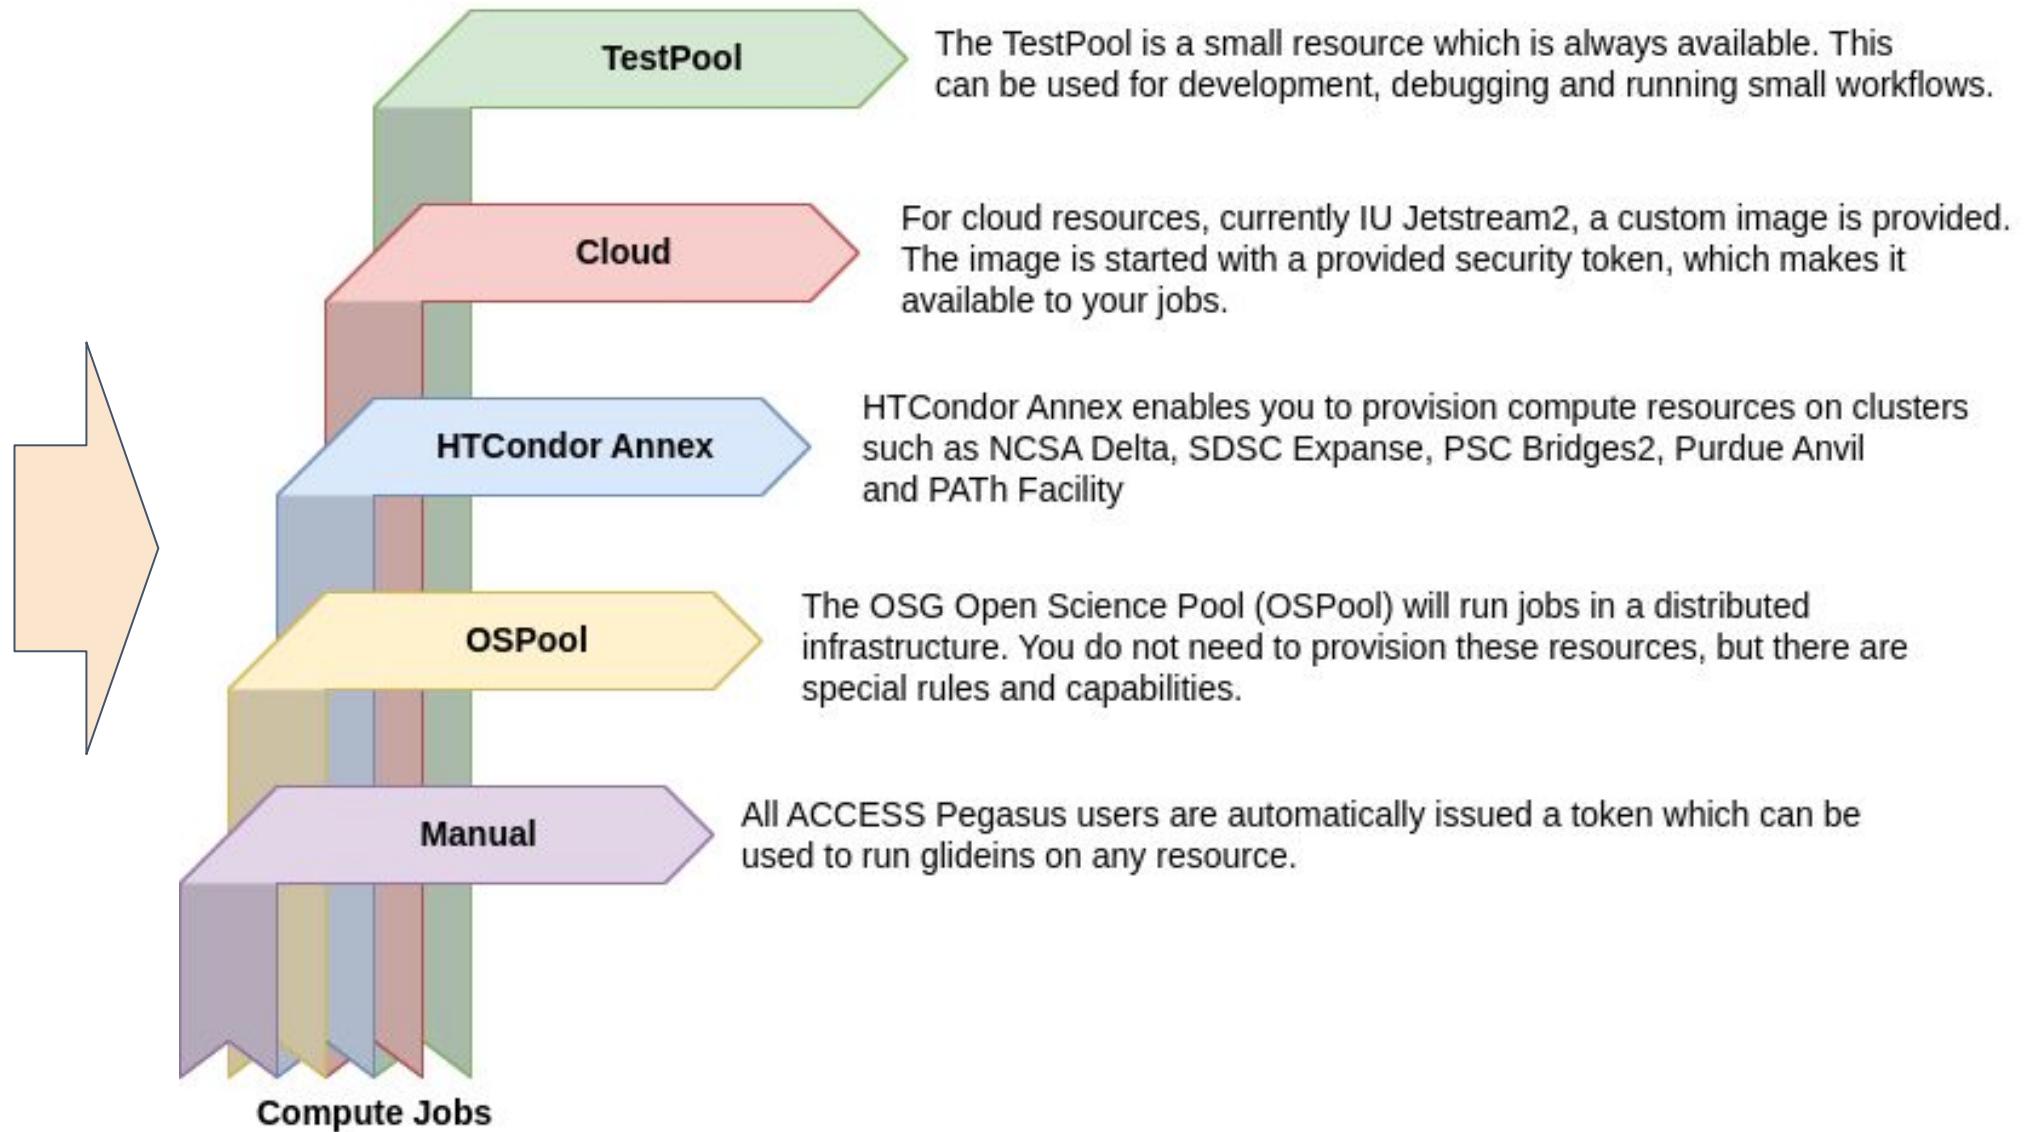
A glidein (pilot) job in HTCondor is a placeholder job that starts on a remote resource, sets up the HTCondor environment, and transforms the resource into a temporary worker node ready to run user jobs from the central HTCondor pool.

Dynamic Resource Provisioning: Jobs start only when needed, reducing idle time.

Late Binding: Matches jobs to the best resources after provisioning.

Central Management: Unified view of all jobs from a single scheduler.

Efficient Setup: Reuses pilots to run multiple jobs, reducing overhead.





TestPool

Helps users with an ACCESS account but no allocation explore the capability

Small amount of compute resources attached to pegasus.access-ci.org

- CPU: 32 cores, 128 GB RAM, 256 GB disk
- GPU: 32 cores, 2 GPUs, 128 GB RAM, 256 GB disk
- Hosted on IU Jetstream2, provisioned when needed

Always available, **no allocation needed**

Can be used for quick turnaround jobs

- workflow development and debugging
- tutorials (not all users might have an allocation at the time of the tutorial)

Cloud

- IU JetStream2
- Provided VM image
- Users have to add pegasus.access-ci.org username and token in the cloud-init yaml
- Instances self-terminates when there are no more jobs

Boot Script

This [cloud-init](#) ⓘ config describes how to provision the instance. It's provided here to permit specific changes in rare circumstances; please modify it cautiously.

 By editing this it's possible to break various Exosphere features like web desktop, web shell, usage graphs, setup status, etc.

```
#cloud-config
users:
  - default
  - name: exouser
    shell: /bin/bash
    groups: sudo, admin
    sudo: ['ALL=(ALL) NOPASSWD:ALL']{ssh-authorized-keys}
  ssh_pwauth: true
  package_update: true
  package_upgrade: {install-os-updates}
  packages:
    - git{write-files}
bootcmd:
  - /opt/ACCESS-Pegasus-Jetstream2/bin/vm-conf alice
  aabbcc...
runcmd:
  - echo on > /proc/sys/kernel/printk_devkmsg || true
  # Disable console rate limiting for distros that use kmsg
  - sleep 1 # Ensures that console log output from any previous command completes before the following command begins
```



delta
cpu
cpu-interactive
gpuA100x4
gpuA100x4-preempt
gpuA100x4-preempt
gpuA100x8
gpuA40x4
gpuA40x4-preempt

stampede2
normal
development
skx-normal

expanser
compute
gpu
shared
gpu-shared

anvil
wholenode
wide
shared
gpu
gpu-debug

bridges2
RM
RM-512
RM-shared
EM
GPU
GPU-shared

path-facility
cpu

HTCondor Annex

- Bring your own HPC allocation
- Semi-managed, submits glideins via SSH.
 - A glidein can run multiple user jobs - it stays active until no more user jobs are available or until end of life has been reached, whichever comes first.
 - A glidein is partitionable - job slots will dynamically be created based on the resource requirements in the user jobs. This means you can fit multiple user jobs on a compute node at the same time.
 - A glidein will only run jobs for the user who started it.
- Documentation: <https://htcondor.org/experimental/ospool/byoc/>

```
$ htcondor annex create --nodes 1 --lifetime 7200 \
--project sta230005p --gpu-type v100-16 $USER GPU@bridges2
```



pegasus-glidein

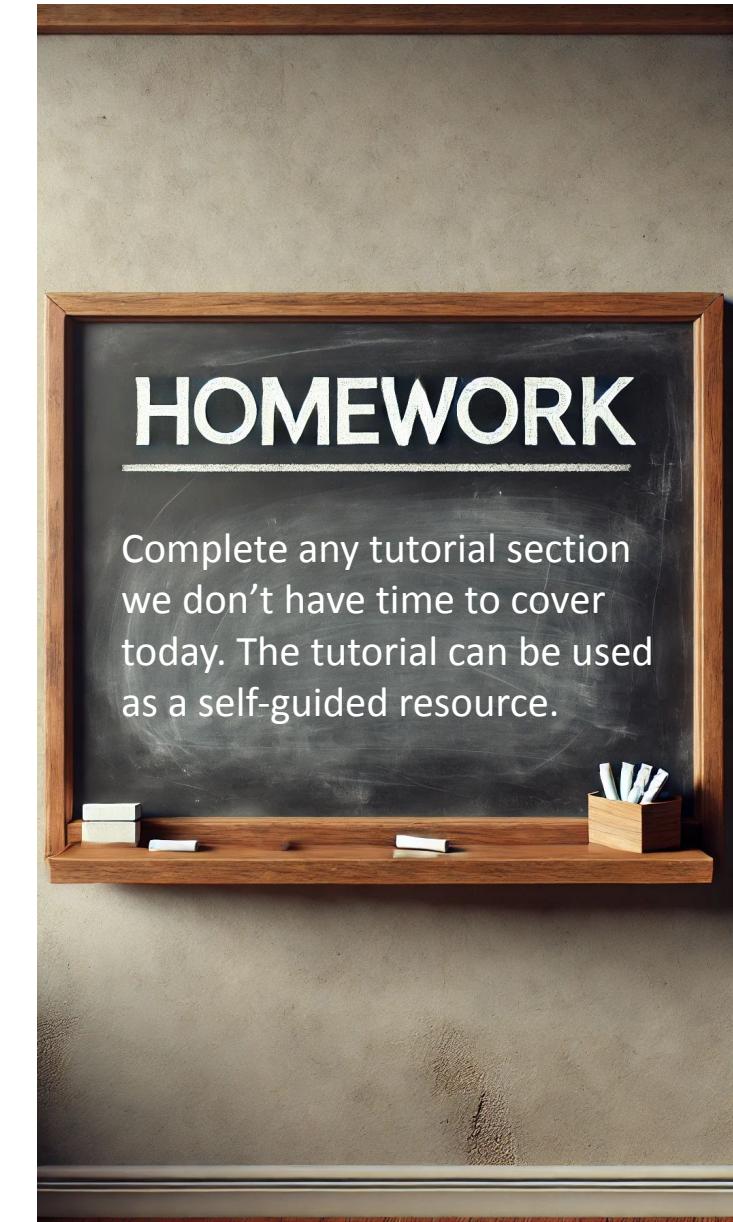
- Simple glidein which can be run anywhere, as long as you have outbound network connectivity
- Ties in well with ACCESS Pegasus
- <https://github.com/pegasus-isi/pegasus-glidein>

```
#!/bin/bash
#SBATCH --job-name=glidein
#SBATCH --nodes=10
#SBATCH --ntasks-per-node=1
#SBATCH --cpus-per-task=48
#SBATCH --time=24:00:00

curl -o pegasus-glidein https://raw.githubusercontent.com/pegasus-isi/pegasus-glidein/main/pegasus-glidein
chmod a+x pegasus-glidein
srun ./pegasus-glidein -c pegasus.access-ci.org \
    -t mytoken \
    -s 'RemoteOwner == "myusername"'
```



- 1. Use a workflow system for your AI workloads**
- 2. Scheduler overlay to simplify resource management**





Thank You!

<https://pegasus.isi.edu>

Pegasus Users Slack and mailing lists: <https://pegasus.isi.edu/contact/>

E-mail Support: pegasus-support@isi.edu

Office Hours: <https://pegasus.isi.edu/office-hours/>